

# SIMPLE EXAMPLES OF THE USE OF MIF WITHIN THE POMP PACKAGE

AARON A. KING

In this vignette, we show the use of the `POMPPackage`, and the `MIF` algorithm within it by means of simple examples. The first is a trivial example: much simpler methods could be used for this model. Nevertheless, it will serve to illustrate the functionality provided by `POMP` and to give some hints as to its use. The second example is nontrivial and involves a model for which alternate fitting procedures are available. One can therefore check the consistency of the results `MIF` gives.

## 1. GETTING STARTED WITH POMP

### 2. A TRIVIAL EXAMPLE: IID NOISE

A trivial example of a state-space model is afforded by the case of an i.i.d. random variable. The example is useful insofar as it allows us to demonstrate the use of the `pomp` package.

**2.1. The model.** We will denote the unobserved (state) process by  $X_t$  and the observed process by  $Y_t$ . To be specific, assume that the unobserved process is just a constant process:  $X_t = \mu$ . The observation process is just measurement with normal error:

$$Y_t \sim N(X_t, \sigma^2).$$

So  $Y_t \sim N(\mu, \sigma^2)$ . Although it isn't natural in this problem, many other problems will have a additional parameters, namely, the initial values of the state variables. Let's assume that we have an initial condition in this problem as well:  $X_0 = x_0$ . We'll have to estimate  $\mu$ ,  $\sigma$ , and  $x_0$ .

**2.2. Inference using `pomp`.** The basic element of the `pomp` package is the `pomp`, or "partially-observed Markov process" object. Let's set up such an object for the IID model above using simulated data. First, we'll need to load the package:

```
> require(pomp)
```

We'll use the following parameters for generating the simulated data.

---

*Date:* May 22, 2007.

```
> true.params <- c(m = 5, s = 3, x0 = 10)
```

Now we come to the setup of the `pomp` object itself. First we'll generate some simulated data:

```
> x <- with(as.list(true.params), rbind(time = seq(1, 100), obs = rnorm(100,
+   mean = m, sd = s)))
```

Next, we'll need four basic functions, which define the model we're interested in. These are

**rprocess:** This function simulates one step of the unobserved process, i.e., it is a random draw from the conditional distribution  $f(X_{t+1} | X_t)$ .

**rmeasure:** This simulates the measurement process, given the state. That is, it is a random draw from the conditional distribution  $f(Y_t | X_t)$ .

**dmeasure:** This furnishes the probability density or mass of the measurement process, i.e., it gives  $f(Y_t = y | X_t)$ .

**particles:**

```
> rprocess <- function(X, t1, t2, ...) {
+   X["x", ] <- X["m", ]
+   X
+ }
> rmeasure <- function(X, time, ...) {
+   rnorm(ncol(X), mean = X["x", ], sd = X["s", ])
+ }
> dmeasure <- function(X, Y, ...) {
+   dnorm(Y["obs"], mean = X["x", ], sd = X["s", ])
+ }
> particles <- function(Np, center, sd, ...) {
+   X <- matrix(0, nrow = length(center) + 1, ncol = Np)
+   rownames(X) <- c("x", names(center))
+   X[names(center), ] <- rnorm(Np * length(center), mean = center,
+     sd = sd)
+   X["x", ] <- X["x0", ]
+   X
+ }
> iid <- pomp(data = x, t0 = 0, rprocess = rprocess, rmeasure = rmeasure,
+   dmeasure = dmeasure, particles = particles)
```

We can simulate from the `pomp` object:

```
> y <- simulate(iid, coef = c(m = 9, s = 1, x0 = 100))
```

We can plot the results (Fig. 1).

```
> plot(y)
```




FIGURE 1. Simulating from an `pomp` object. The command used to generate the plot is shown above.

Now we can use the MIF algorithm to find maximum likelihood estimates of the parameters.

```
> z <- mif(iid, start = c(m = 3, s = 5, x0 = 10), rw.sd = c(m = 0.1,
+   s = 0.1, x0 = 1), Nmif = 1, ivps = "x0", pars = c("m", "s"),
+   stvs = "x", alg.pars = list(cooling.factor = 0.99, CC = 2,
+   Np = 50, T0 = 3))
> z <- continue(z, Nmif = 50, alg.pars = list(cooling.factor = 0.99,
+   CC = 5, Np = 50, T0 = 3))
```

To perform additional MIF iterations on a `mif` object, one uses the `continue` function.

```
> z <- continue(z, Nmif = 500)
> truth <- c(true.params, loglik = sum(dnorm(iid@data[2, ], mean = 5,
+   sd = 3, log = T)))
> mle.params <- c(m = mean(iid@data[2, ]), s = sd(iid@data[2, ]),
+   x0 = NA)
```

```
> predvarplot(z)
```




FIGURE 2. Using `predvarplot` to choose a good value of CC. The command used to generate the plot is shown above. A value of 5 looks good.

```
> mle <- c(mle.params, loglik = with(as.list(mle.params), sum(dnorm(iid@data[
+   ], mean = m, sd = s, log = T))))
> best <- pfilter(z)
> best.estimate <- c(coef(z), best$loglik)
> require(xtable)
> xtable(cbind(truth = truth, mle = mle, "MIF estimate" = best.estimate),
+   caption = "Comparison of true, MLE, and best MIF parameters",
+   digits = c(0, 3, 3, 3))
```

```
> par(mfrow = c(2, 4))
> plot(z)
> par(mfrow = c(1, 1))
```

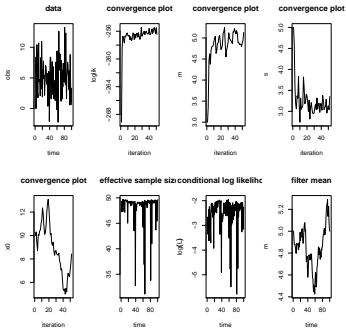


FIGURE 3. Plotting a mif object.

	truth	mle	MIF estimate
m	5.000	4.978	5.106
s	3.000	3.085	3.033
x0	10.000		5.498
loglik	-254.119	-254.065	-254.168

TABLE 1. Comparison of true, MLE, and best MIF parameters

### 3. A NONTRIVIAL EXAMPLE: A 2-D ORNSTEIN-UHLENBECK PROCESS

We will want to maximize the likelihood on a transformed parameter space. It is useful to define functions that will perform the forward and inverse parameter transformations.

```
> par(mfrow = c(2, 4))
> plot(z)
> par(mfrow = c(1, 1))
```

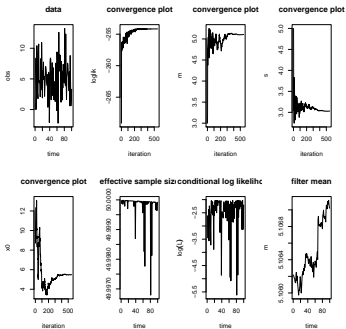


FIGURE 4. Results of fitting the model to the data using MIF.

```
> par.trans <- function(params) {
+   x <- params
+   r <- length(dim(x))
+   if (r > 1) {
+     x <- apply(x, 2:r, par.trans)
+   }
+   else {
+     x["alpha.1"] <- log(-params["alpha.1"])
+     x["alpha.2"] <- log(params["alpha.2"])
+     x["alpha.3"] <- log(-params["alpha.3"])
+     x["alpha.4"] <- log(-params["alpha.4"])
+     x["sigma.1"] <- log(params["sigma.1"])
+     x["sigma.3"] <- log(params["sigma.3"])
+     x["tau"] <- log(params["tau"])
+   }
+ }
```

```

+   x
+ }
> par.untrans <- function(params) {
+   x <- params
+   r <- length(dim(x))
+   if (r > 1) {
+     x <- apply(x, 2:r, par.trans)
+   }
+   else {
+     x["alpha.1"] <- -exp(params["alpha.1"])
+     x["alpha.2"] <- exp(params["alpha.2"])
+     x["alpha.3"] <- -exp(params["alpha.3"])
+     x["alpha.4"] <- -exp(params["alpha.4"])
+     x["sigma.1"] <- exp(params["sigma.1"])
+     x["sigma.3"] <- exp(params["sigma.3"])
+     x["tau"] <- exp(params["tau"])
+   }
+   x
+ }

> rw.sd <- c(alpha.1 = 0.1, alpha.2 = 0.1, alpha.3 = 0.1, alpha.4 = 0.1,
+   sigma.1 = 0.1, sigma.2 = 0.1, sigma.3 = 0.1, tau = 0.1, x1.0 = 0.2,
+   x2.0 = 0.2)
> start <- c(alpha.1 = -0.1, alpha.2 = 0.1, alpha.3 = -0.2, alpha.4 = -0.5,
+   sigma.1 = 0.1, sigma.2 = 0, sigma.3 = 2, tau = 2, x1.0 = 5,
+   x2.0 = 3)
> truth <- c(alpha.1 = -0.5, alpha.2 = 0.1, alpha.3 = -0.1, alpha.4 = -0.5,
+   sigma.1 = 1, sigma.2 = 0, sigma.3 = 1, tau = 0.5, x1.0 = 5,
+   x2.0 = 3)
> estnames <- c("alpha.1", "alpha.3", "sigma.1", "sigma.3", "tau")
> fixnames <- c("alpha.2", "alpha.4", "sigma.2")
> ivpnames <- c("x1.0", "x2.0")

> rprocess <- function(X, t1, t2, ...) {
+   nvar <- nrow(X)
+   np <- ncol(X)
+   stateindex <- match(c("x1"), rownames(X)) - 1
+   parindex <- match(c("alpha.1", "alpha.2", "alpha.3", "alpha.4",
+     "sigma.1", "sigma.2", "sigma.3"), rownames(X)) - 1
+   result <- .C("ou2_adv", X = as.double(X), nvar = as.integer(nvar),
+     np = as.integer(np), stateindex = as.integer(stateindex),
+     parindex = as.integer(parindex), DUP = FALSE, NAOK = TRUE,
+     PACKAGE = "pomp")$X

```

```

+   array(result, dim = dim(X), dimnames = dimnames(X))
+ }
> dmeasure <- function(X, Y, ...) {
+   n <- dim(X)
+   measindex <- match(c("x1", "tau"), rownames(X)) - 1
+   .C("normal_dmeasure", n = as.integer(n), index = as.integer(measindex),
+     X = as.double(X), y = as.double(Y[2]), f = double(n[2]),
+     DUP = FALSE, NAOK = TRUE, PACKAGE = "pomp")$f
+ }
> rmeasure <- function(X, time, ...) {
+   n <- dim(X)
+   nsim <- n[2]
+   measindex <- match(c("x1", "tau"), rownames(X)) - 1
+   matrix(.C("normal_rmeasure", n = as.integer(n), index = as.integer(measindex),
+     X = as.double(X), cases = double(nsim), DUP = FALSE,
+     NAOK = TRUE, PACKAGE = "pomp")$cases, 1, nsim)
+ }
> particles <- function(Np, center, sd, statenames, parnames, fixnames,
+   ivpnames, ...) {
+   X <- matrix(data = 0, nrow = length(statenames) + length(parnames) +
+     length(fixnames) + length(ivpnames), ncol = Np, dimnames = list(c(statenames,
+     parnames, fixnames, ivpnames), NULL))
+   X[parnames, ] <- rnorm(n = Np * length(parnames), mean = center[parnames],
+     sd = sd[parnames])
+   X[fixnames, ] <- center[fixnames]
+   X[ivpnames, ] <- rnorm(n = Np * length(ivpnames), mean = center[ivpnames],
+     sd = sd[ivpnames])
+   X[statenames, ] <- X[ivpnames, ]
+   X
+ }
> ex <- pomp(data = rbind(t = 1:1000, z = rep(0, 1000)), t0 = 0,
+   rprocess = rprocess, dmeasure = dmeasure, rmeasure = rmeasure,
+   particles = particles, parnames = estnames, fixnames = fixnames,
+   ivpnames = c("x1.0", "x2.0"), statenames = c("x1", "x2"))
> ex <- simulate(ex, coef = par.trans(truth))
> print(arima.fit <- arima(ex@data[2, ], order = c(2, 0, 2)))

Call:
arima(x = ex@data[2, ], order = c(2, 0, 2))

Coefficients:
      ar1      ar2      ma1      ma2  intercept

```




```

      -0.0871  0.1424  -0.3373  0.0650   -0.0221
s.e.   0.3767  0.2091   0.3765  0.0836   0.0273

sigma^2 estimated as 1.258: log likelihood = -1533.84, aic = 3079.68
> ff <- pfilter(ex, Np = 1000, coef = par.trans(truth))
> print(c(loglik = ff$loglik, nfail = ff$nfail))
      loglik      nfail
-1535.221      0.000

> x <- mif(ex, start = par.trans(start), rw.sd = rw.sd, Nmif = 1,
+         pars = estnames, ivps = ivpnames, stvs = c("x1", "x2"), alg.pars = list(
+           CC = 7, T0 = 4, cooling.factor = 0.95), max.fail = 1000,
+         weighted = F)
> predvarplot(x)

```



```

> x <- continue(x, Nmif = 20, max.fail = 100, weighted = F)
> x <- continue(x, Nmif = 80, max.fail = 100, weighted = T)

> plot(x)
> print(x)

1 -variable time series of length: 1000
number of MIF iterations done: 101

```

5 parameter(s) estimated: alpha.1, alpha.3, sigma.1, sigma.3, tau  
 2 state variable(s) (STV): x1, x2  
 2 initial-value parameter(s) (IVP): x1.0, x2.0

Sigma (parameter scales):

alpha.1	alpha.2	alpha.3	alpha.4	sigma.1	sigma.2	sigma.3	tau	x1.0	x2.0
0.1	0.1	0.1	0.1	0.1	0.1	0.1	0.1	0.2	0.2

number of particles = 1000

CC = 7 , T0 = 4

cooling factor = 0.95

random walk intensity at last MIF iteration:


alpha.1	alpha.2	alpha.3	alpha.4	sigma.1	sigma.2	sigma.3	tau
0.000592	0.000592	0.000592	0.000592	0.000592	0.000592	0.000592	0.000592
	x1.0	x2.0					
0.001184	0.001184						

estimated parameter(s):

alpha.1	alpha.3	sigma.1	sigma.3	tau	x1.0	x2.0
-1.290	-1.385	-1.109	1.071	-0.365	6.303	5.475

log-likelihood (w/ variable parameters): -1537

number of filtering failures: 0



```
> ff.fit <- pfilter(x)
> print(c(loglik = ff.fit$loglik, nfail = ff.fit$nfail))
```

```
      loglik      nfail
-1534.519      0.000
```

```
> print(c(loglik = ff$loglik, nfail = ff$nfail))
```

```
      loglik      nfail
-1535.221      0.000
```

```
> print(arima.fit)
```

Call:

```
arima(x = ex@data[2, ], order = c(2, 0, 2))
```

Coefficients:

	ar1	ar2	ma1	ma2	intercept
	-0.0871	0.1424	-0.3373	0.0650	-0.0221
s.e.	0.3767	0.2091	0.3765	0.0836	0.0273

sigma^2 estimated as 1.258: log likelihood = -1533.84, aic = 3079.68

A. A. KING, DEPARTMENTS OF ECOLOGY & EVOLUTIONARY BIOLOGY AND  
MATHEMATICS, UNIVERSITY OF MICHIGAN, ANN ARBOR, MICHIGAN 48109-1048  
USA

*E-mail address:* [aaron.king@umich.edu](mailto:aaron.king@umich.edu)

*URL:* <http://www.umich.edu/~kingaa>